

In [1]:

```
import pandas as pa
import numpy as np
import scipy
import statsmodels.tsa
import matplotlib.pyplot as plt
```

In [2]:

```
data = pa.read_csv(r'C:\Users\SACHIN K M\Desktop\python\data\datasets\LYNXdata.csv', header=0, names = ['year', 'trappings'], index_col = 0)
```

In [3]:

```
data.head()
```

Out[3]:

trappings	
year	
1821	269
1822	321
1823	585
1824	871
1825	1475

In [4]:

```
type(data)
```

Out[4]:

pandas.core.frame.DataFrame

In [5]:

```
data2 = data['trappings']
data2.head()
```

Out[5]:

```
year
1821    269
1822    321
1823    585
1824    871
1825   1475
Name: trappings, dtype: int64
```

In [6]:

```
#converting data frame into series object
#frequency aliases
#A - for year end
#B - business days
#D - day
#H - hour
#T - minutes
#S - seconds

dataseries = pa.Series(data['trappings'].values, index= pa.date_range('31/12/1821', periods = 114, freq = 'A-DEC'))
dataseries.head()
```

Out[6]:

```
1821-12-31    269
1822-12-31    321
1823-12-31    585
```

1824-12-31 871
1825-12-31 1475
Freq: A-DEC, dtype: int64

In [7]:

```
#dicky fuller test
# p value should be less than 0.05 for observation to be stationary

def stationarity_test(timeseries):
    """augumented dicky fuller
    test for statinarity"""
    from statsmodels.tsa.stattools import adfuller
    print ("results of dicky-fuller test:")
    df_test = adfuller(timeseries, autolag = 'AIC')
    df_output = pd.Series(df_test[0:4],
        index = ["test statistics", "p value", "#lags used", "number of observations used"])
    print (df_output)
```

In [8]:

```
stationarity_test(dataseries)
```

results of dicky-fuller test:

test statistics	-2.996304
p value	0.035241
#lags used	7.000000
number of observations used	106.000000

dtype: float64

In [9]:

```
#np.random.normal(1, 3, 300) , 1 is mean, 3 standard deviation, 300 - no of observations generated randomly
stationarity_test(np.random.normal(1, 3, 300))
plt.figure(figsize=(20,7))
plt.plot(np.random.normal(1, 3, 300))
```


results of dicky-fuller test:

test statistics	-1.868927e+01
p value	2.039774e-30
#lags used	0.000000e+00
number of observations used	2.990000e+02

dtype: float64

Out[9]:

[<matplotlib.lines.Line2D at 0x13ebbca4b70>]



In [10]:

```
import statsmodels.graphics
from statsmodels.graphics import tsaplots
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

In [11]:

```
%matplotlib inline
```


```
fig = plt.figure(figsize=(12,7))
```

```
ax1 = fig.add_subplot(211)
```

```
fig = plot_acf(dataseries, lags=20, ax = ax1)
```

```
ax2 = fig.add_subplot(212)
```

```
fig = plot_pacf(dataseries, lags=20, ax = ax2)
```



In [12]:

```
plt.figure(figsize=(12,7))
```

```
dataseries.plot()
```

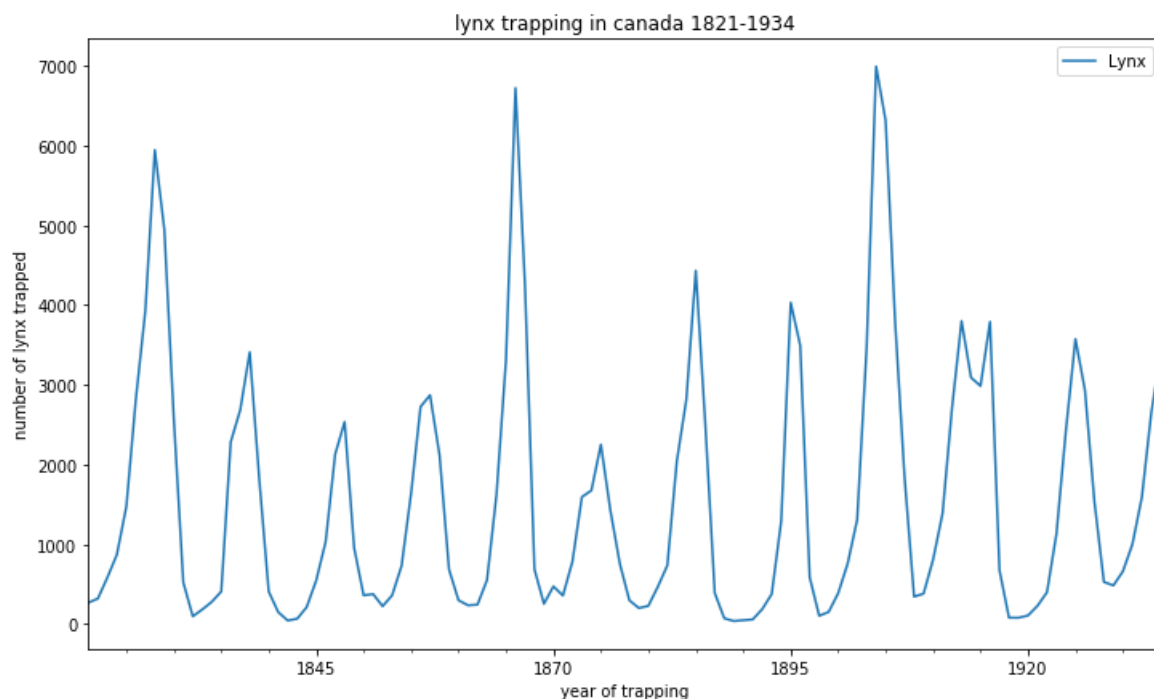
```
plt.title("lynx trapping in canada 1821-1934")
```

```
plt.xlabel('year of trapping')
```

```
plt.ylabel('number of lynx trapped')
```

```
plt.legend(['Lynx'])
```

```
cumsum_lynx = np.cumsum(dataseries)
```



In [13]:

```
cumsum_lynx = np.cumsum(dataseries)
```


```
plt.figure(figsize=(12,7))
```

```
plt.plot(dataseries)
```

```
plt.plot(cumsum_lynx)
plt.title("lynx trapping in canada 1821-1934")
plt.xlabel('year of trapping')
plt.ylabel('number of lynx trapped')
plt.legend(['Lynx per year', 'cummulative total'])
```

Out[13]:

<matplotlib.legend.Legend at 0x13ebbf78d0>



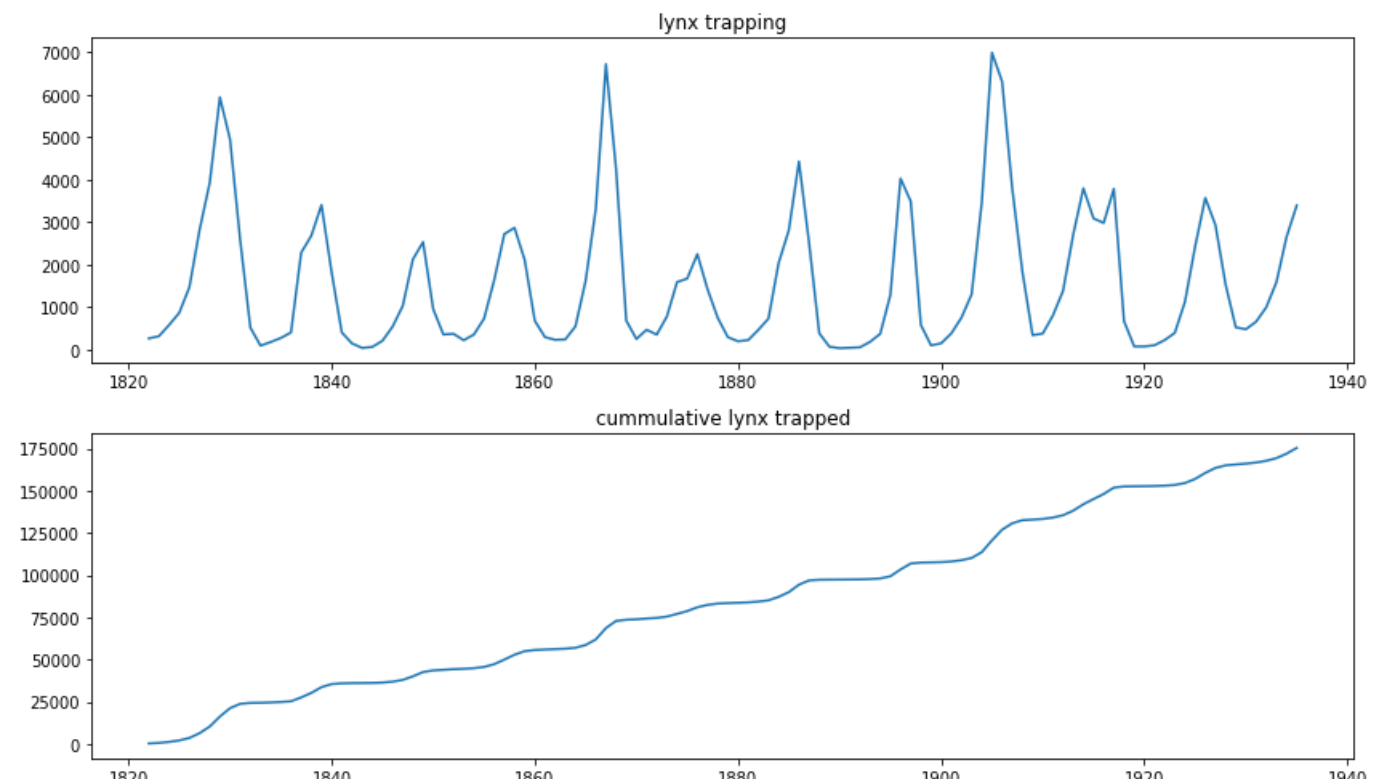
In [14]:

```
plt.figure(figsize=(12,7))

plt.subplot(2, 1, 1)
plt.plot(dataseries)
plt.title('lynx trapping')

plt.subplot(2, 1, 2)
plt.plot(cumsum_lynx)
plt.title('cummulative lynx trapped')

plt.tight_layout()
```



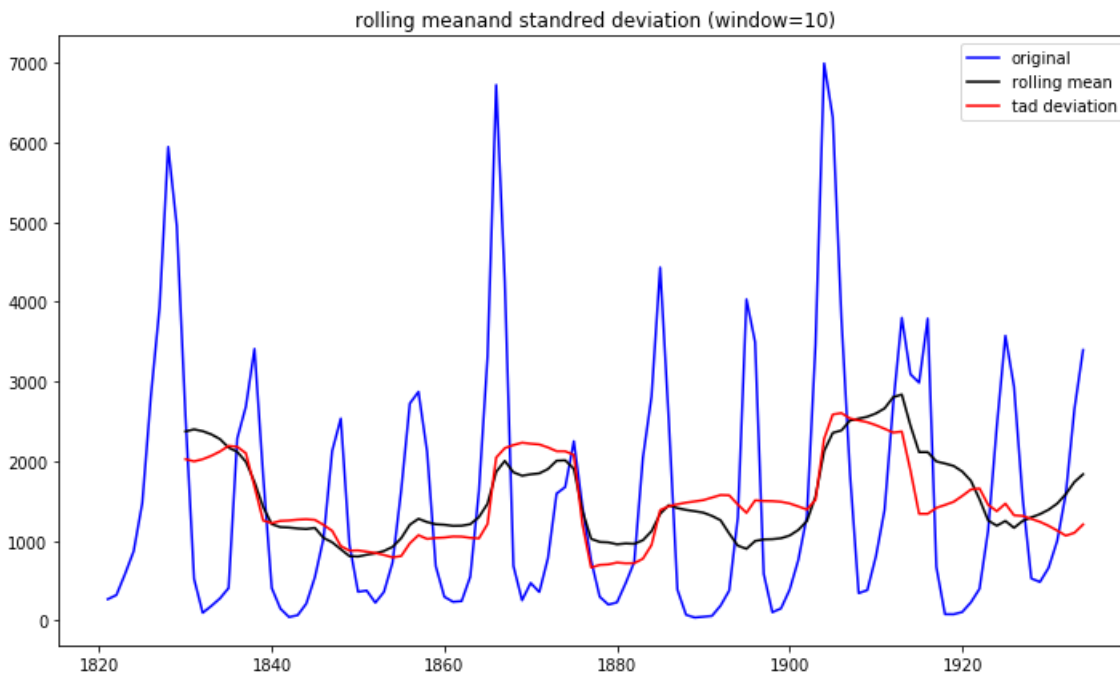
In [15]:

```
def plot_rolling(timeseries, window):
    rol_mean = timeseries.rolling(window).mean()
    rol_std = timeseries.rolling(window).std()

    plt.figure(figsize=(12,7))
    og = plt.plot(timeseries, color="blue", label="original")
    mean = plt.plot(rol_mean, color="black", label="rolling mean")
    std = plt.plot(rol_std, color="red", label="tad deviation")
    plt.legend(loc='best')
    plt.title("rolling meanand standred deviation (window="+str(window)+")")
    plt.show()
```

In [16]:

```
plot_rolling(data, 10)
```



In [17]:

```
#getting smoothing values
dataseries.rolling(10).mean()
dataseries.rolling(10, min_periods=1).mean()
```

Out[17]:

```
1821-12-31    269.000000
1822-12-31    295.000000
1823-12-31    391.666667
1824-12-31    511.500000
1825-12-31    704.200000
1826-12-31   1057.000000
1827-12-31   1467.142857
1828-12-31   2026.625000
1829-12-31   2351.444444
1830-12-31   2374.000000
1831-12-31   2399.400000
1832-12-31   2377.100000
1833-12-31   2337.000000
1834-12-31   2277.800000
1835-12-31   2171.200000
1836-12-31   2117.600000
1837-12-31   1993.300000
1838-12-31   1739.900000
1839-12-31   1427.300000
1840-12-31   1210.500000
1841-12-31   1173.300000
1842-12-31   1168.000000
1843-12-31   1156.400000
1844-12-31   1149.800000
```

```
1845-12-31 1163.500000
1846-12-31 1038.300000
1847-12-31 982.700000
1848-12-31 895.400000
1849-12-31 808.700000
1850-12-31 803.900000
...
```

```
1905-12-31 2356.100000
1906-12-31 2386.000000
1907-12-31 2510.900000
1908-12-31 2534.900000
1909-12-31 2557.800000
1910-12-31 2599.900000
1911-12-31 2662.900000
1912-12-31 2803.500000
1913-12-31 2837.000000
1914-12-31 2447.000000
1915-12-31 2114.200000
1916-12-31 2113.800000
1917-12-31 1997.600000
1918-12-31 1971.200000
1919-12-31 1941.000000
1920-12-31 1871.000000
1921-12-31 1755.100000
1922-12-31 1523.700000
1923-12-31 1256.900000
1924-12-31 1191.000000
1925-12-31 1249.900000
1926-12-31 1164.400000
1927-12-31 1250.700000
1928-12-31 1295.500000
1929-12-31 1336.000000
1930-12-31 1391.400000
1931-12-31 1468.500000
1932-12-31 1587.600000
1933-12-31 1740.100000
1934-12-31 1836.500000
```

Freq: A-DEC, Length: 114, dtype: float64

In [18]:

```
# to avoid nan values at the wbeganing window size
dataseries.head()

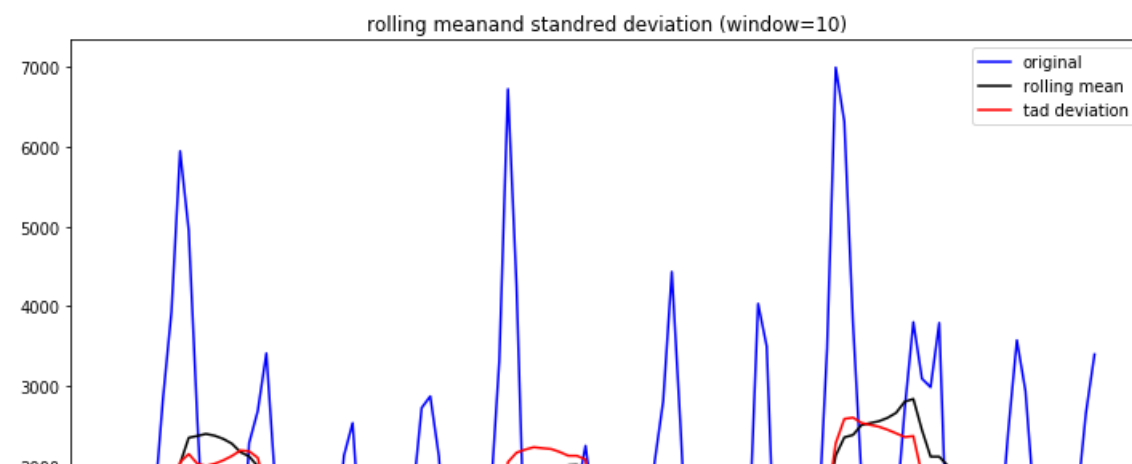
# difference between plot_rolling & plot_rolling1 is min_period it will start from there


def plot_rolling1(timeseries, window):
    rol_mean = timeseries.rolling(window, min_periods=1).mean()
    rol_std = timeseries.rolling(window, min_periods=1).std()

    plt.figure(figsize=(12,7))
    og = plt.plot(timeseries, color="blue", label="original")
    mean = plt.plot(rol_mean, color="black", label="rolling mean")
    std = plt.plot(rol_std, color="red", label="tad deviation")
    plt.legend(loc='best')
    plt.title("rolling meanand standred deviation (window="+str(window)+")")
    plt.show()
```

In [19]:

```
plot_rolling1(dataseries, 10)
```





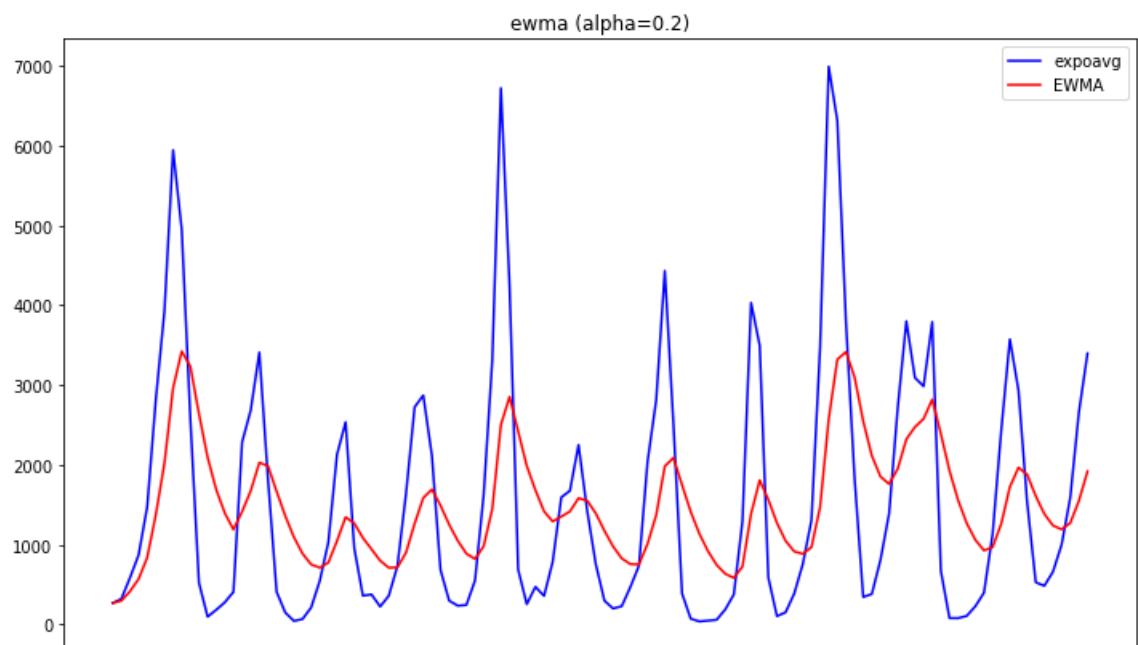
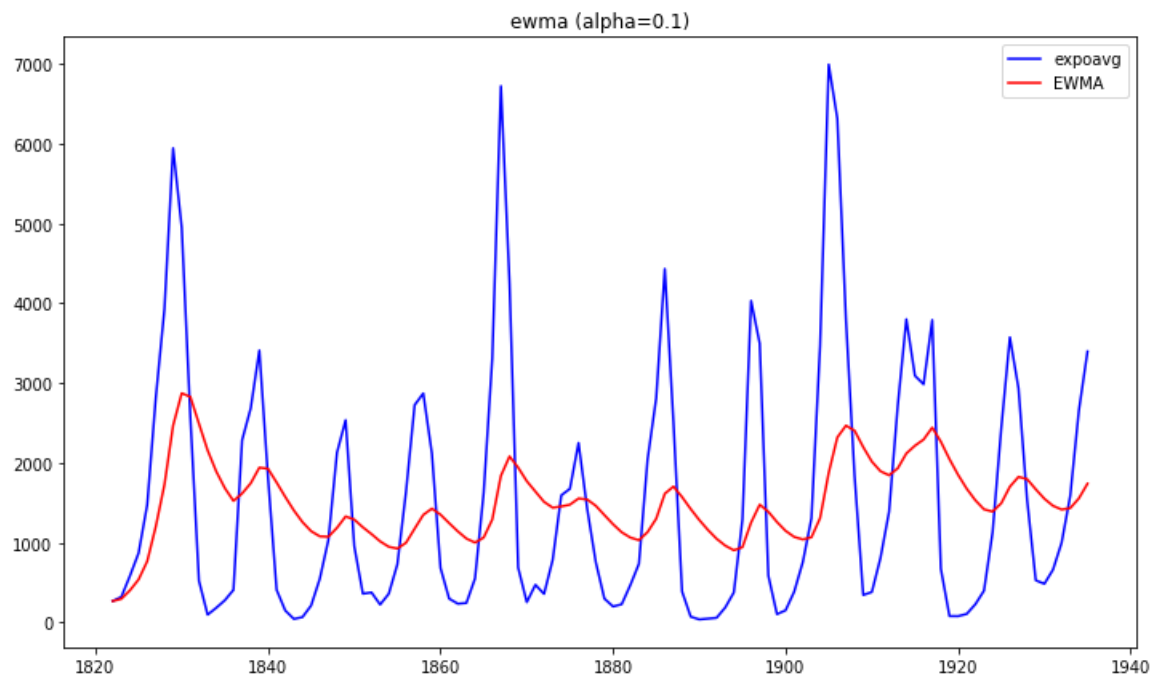
In [20]:


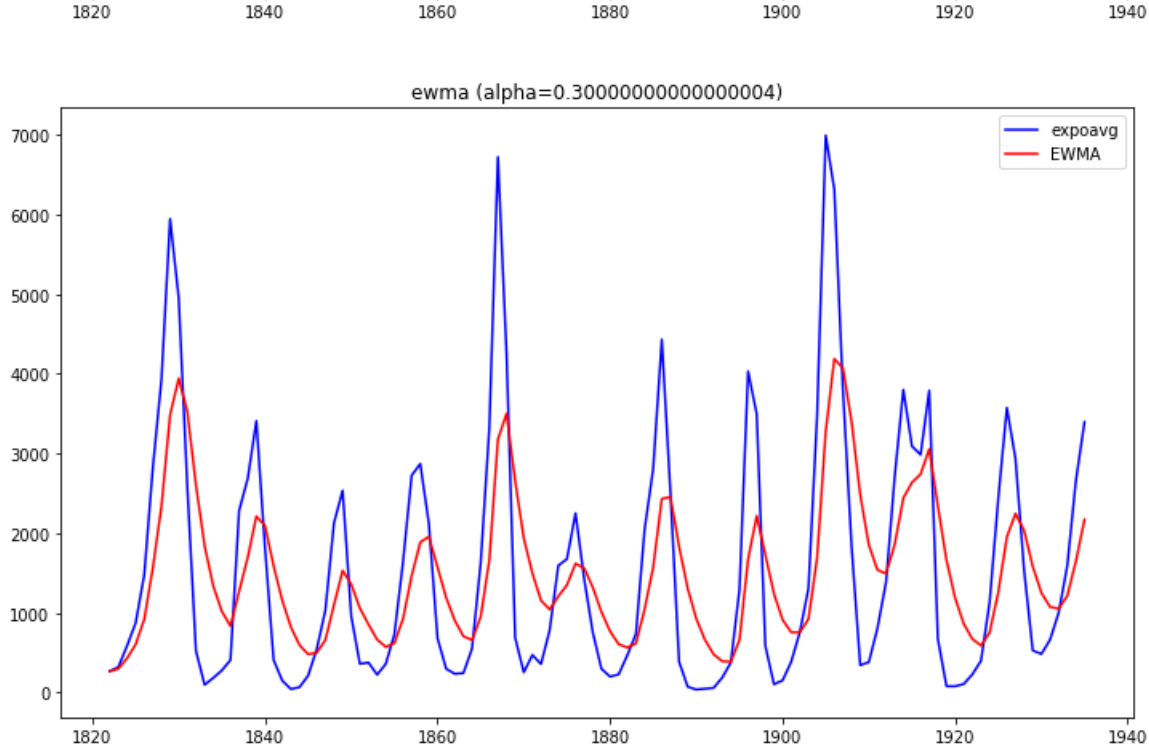
```
#ewma (exponential weighted moving average)  
#this method only for pandas data series or pandas data frame
```

```
def plot_ewma(timeseries, alpha):  
    expw_ma = timeseries.ewm(alpha=alpha).mean()  
  
    fig= plt.figure(figsize=(12,7))  
    og_line = plt.plot(timeseries, color= 'blue', label='expoavg')  
    exwm_avg = plt.plot(expw_ma, color='red', label='EWMA')  
    plt.legend(loc='best')  
    plt.title("ewma (alpha="+str(alpha)+"")  
    plt.show()
```

In [21]:

```
for i in np.arange(0.1, 0.5, 0.1):  
    plot_ewma(dataseries, i)
```





In [22]:


```
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import acf, pacf
from statsmodels.tsa.arima_model import ARIMA
```


In [23]:

```
model = ARIMA(dataseries, order=(2, 0, 0))
results_AR = model.fit()
plt.figure(figsize=(20,7))
plt.plot(dataseries)
plt.plot(results_AR.fittedvalues, color='red')
```

Out[23]:

[<matplotlib.lines.Line2D at 0x13ebc351908>]





In [24]:

```
#we will consider AIC value it should be as low as possible
#only randomness present in residual
results_AR.summary()
```

Out[24]:

ARMA Model Results

Dep. Variable:	y	No. Observations:	114
Model:	ARMA(2, 0)	Log Likelihood	-935.016
Method:	css-mle	S.D. of innovations	876.447
Date:	Fri, 12 Jul 2019	AIC	1878.032
Time:	07:43:12	BIC	1888.977
Sample:	12-31-1821	HQIC	1882.474
	- 12-31-1934		


	coef	std err	z	P> z	[0.025	0.975]
const	1545.3385	181.671	8.506	0.000	1189.269	1901.408
ar.L1.y	1.1474	0.074	15.459	0.000	1.002	1.293
ar.L2.y	-0.5997	0.074	-8.110	0.000	-0.745	-0.455

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	0.9566	-0.8673j	1.2913	-0.1172
AR.2	0.9566	+0.8673j	1.2913	0.1172

In [25]:

```
#ACF on our residual model
fig = plt.figure(figsize=(12,7))
ax1 = fig.add_subplot(211)
fig = plot_acf(results_AR.resid, lags=20, ax= ax1)
```



In [26]:

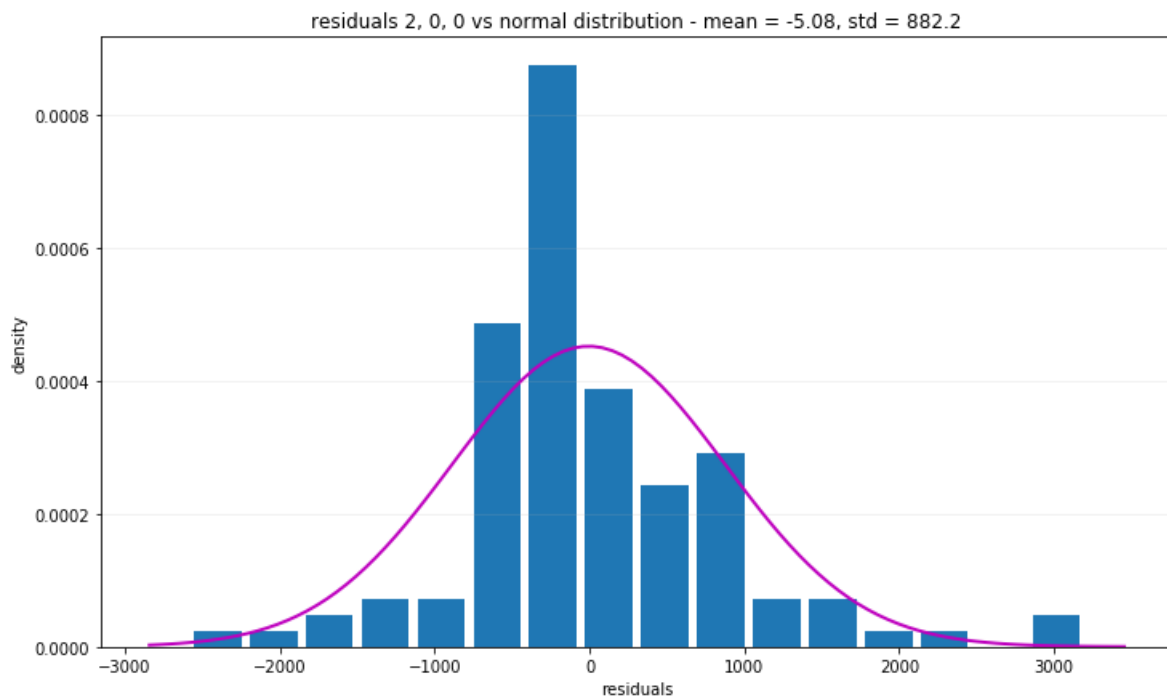
```
from scipy.stats import norm
```

```
plt.figure(figsize=(12,7))
```

```

plt.figure(figsize=(12,7))
plt.hist(results_AR.resid, bins = 'auto', density = True, rwidth = 0.85,
        label = 'residual')
mu, std = norm.fit(results_AR.resid)
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = norm.pdf(x, mu, std)
plt.plot(x, p, 'm', linewidth = 2)
plt.grid(axis = 'y', alpha = 0.2)
plt.xlabel('residuals')
plt.ylabel('density')
plt.title('residuals 2, 0, 0 vs normal distribution - mean = '+str(round(mu, 2))+', std = '+str(round(std, 2)))
plt.show()

```



In [27]:

#order=(4, 0, 0) is producing the lowest value of AIC AR5 & AR3 both are higher value of AIC


```

model = ARIMA(dataseries, order=(4, 0, 0))
results_AR1 = model.fit()
plt.figure(figsize=(20,7))
plt.plot(dataseries)
plt.plot(results_AR.fittedvalues, color='red')

```

Out[27]:

[<matplotlib.lines.Line2D at 0x13ebbef25c0>]



In [28]:

```

results_AR1.summary()

```

Out[28]:

ARMA Model Results

Dep. Variable:	y	No. Observations:	114
Model:	ARMA(4, 0)	Log Likelihood	-931.111
Method:	css-mle	S.D. of innovations	845.949
Date:	Fri, 12 Jul 2019	AIC	1874.222
Time:	07:45:46	BIC	1890.639
Sample:	12-31-1821	HQIC	1880.885
	- 12-31-1934		

	coef	std err	z	P> z	[0.025	0.975]
const	1547.4367	136.851	11.307	0.000	1279.214	1815.659
ar.L1.y	1.1246	0.090	12.450	0.000	0.948	1.302
ar.L2.y	-0.7174	0.137	-5.250	0.000	-0.985	-0.450
ar.L3.y	0.2634	0.136	1.935	0.056	-0.003	0.530
ar.L4.y	-0.2543	0.090	-2.837	0.005	-0.430	-0.079

Roots

	Real	Imaginary	Modulus	Frequency
AR.1	0.9198	-0.6880j	1.1486	-0.1022
AR.2	0.9198	+0.6880j	1.1486	0.1022
AR.3	-0.4020	-1.6789j	1.7264	-0.2874
AR.4	-0.4020	+1.6789j	1.7264	0.2874

In [29]:

```
results_AR1.resid.tail()
```

Out[29]:

1930-12-31 -65.572508
1931-12-31 -48.257955
1932-12-31 43.827806
1933-12-31 631.973963
1934-12-31 550.263041
Freq: A-DEC, dtype: float64

In [30]:

```
results_AR1.fittedvalues.tail()
```

Out[30]:

1930-12-31 727.572508
1931-12-31 1048.257955
1932-12-31 1546.172194
1933-12-31 2025.026037
1934-12-31 2845.736959
Freq: A-DEC, dtype: float64

In [31]:

```
dataseries.tail()
```

Out[31]:

1930-12-31 662
1931-12-31 1000
1932-12-31 1590
1933-12-31 2657
1934-12-31 3396
Freq: A-DEC, dtype: int64

In [32]:


```
np.mean(results_AR1.resid)
```

Out[32]:

-9.065780174278016

In [33]:

```
fig = plt.figure(figsize=(12,7))
ax1 = fig.add_subplot(211)
fig = plot_acf(results_AR1.resid, lags=20, ax= ax1)
```

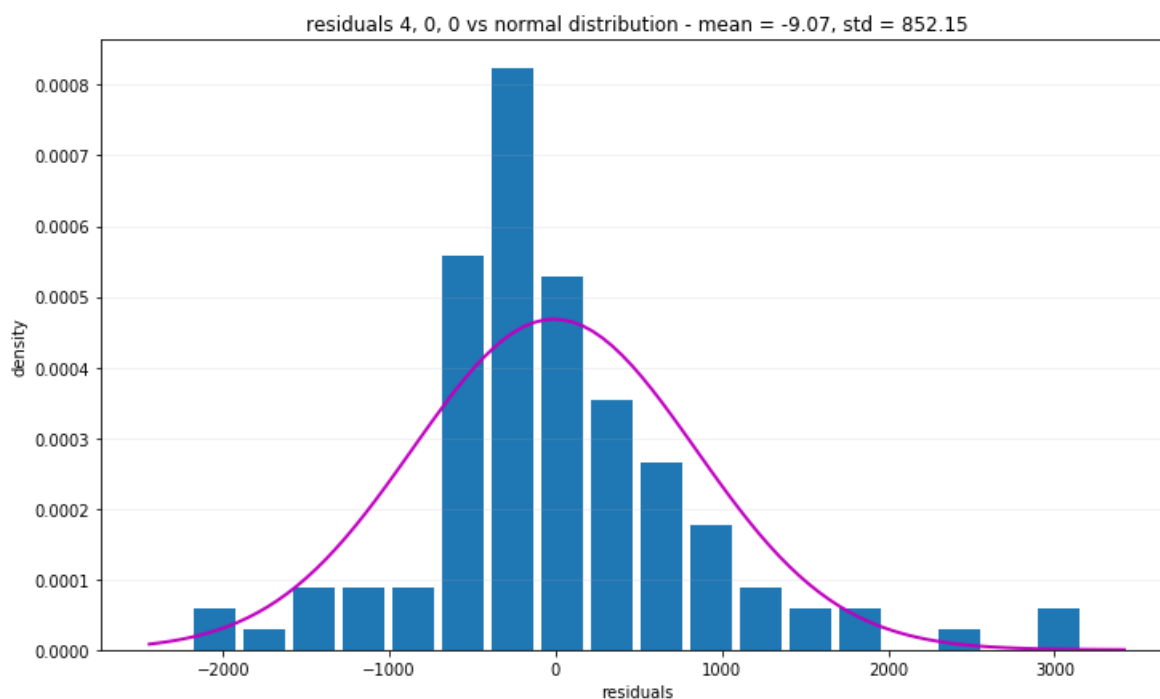


In [35]:

```
#mu is mean, pdf --> probability density function
#line space is used for to equally space the bars in graph
from scipy.stats import norm

plt.figure(figsize=(12,7))
plt.hist(results_AR1.resid, bins = 'auto', density = True, rwidth = 0.85,
         label = 'residual')
mu, std = norm.fit(results_AR1.resid)
print(mu, std)
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = norm.pdf(x, mu, std)
plt.plot(x, p, 'm', linewidth = 2)
plt.grid(axis = 'y', alpha = 0.2)
plt.xlabel('residuals')
plt.ylabel('density')
plt.title('residuals 4, 0, 0 vs normal distribution - mean = '+str(round(mu, 2))+', std = '+str(round(std, 2)))
plt.show()
```

-9.065780174278022 852.1504376527944



In [36]:

```
model202 = ARIMA(dataseries, order=(2, 0, 2))
results_AR2 = model202.fit()
```

In [38]:

```
Fcast400 = results_AR1.predict(start = '31/12/1935', end = '31/12/1945')
fcast202 = results_AR2.predict(start = '31/12/1935', end = '31/12/1945')
```


C:\Users\SACHIN K M\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:336: FutureWarning: Creating a DatetimeIndex by passing range endpoints is deprecated. Use `pandas.date_range` instead.
freq=base_index.freq)

In [40]:

```
plt.figure(figsize=(12,7))
plt.plot(dataseries, linewidth=2, label = "original")
plt.plot(Fcast400, color='red', linewidth=2, label="ARIMA 4, 0, 0")
plt.plot(fcast202, color='blue', linewidth=2, label="ARIMA 2, 0, 2")
plt.legend()
```

Out[40]:

<matplotlib.legend.Legend at 0x13ebccce710>



In []: